

Guardrails-as-a-Service: Hybrid Deterministic and Semantic Policy Evaluation

Pradeep Annepu, Rajesh Akula
Email: {M25AI1109, M25AI1109}@iitj.ac.in

Abstract—We present the architecture and planned evaluation of a cloud-agnostic Guardrails-as-a-Service platform that integrates deterministic rule evaluation with semantic vector retrieval to enforce complex governance policies for AI-enabled systems. The design emphasizes scalability via stateless microservices and asynchronous event flow, extensibility through a handler abstraction, and audit integrity using hash-chained decision logs. We define measurable quality attributes (latency, throughput, extensibility effort, audit lag) and outline empirical validation strategies supporting operational assurance and trust.

Index Terms—Policy enforcement, microservices, vector similarity, audit logging, hash chaining, observability, cloud governance.

I. INTRODUCTION

Dynamic policy enforcement in modern distributed and AI-centric systems requires synthesizing explicit rules with semantic context. Traditional engines (e.g., Open Policy Agent (OPA) [1]) optimize declarative logic but lack native semantic retrieval. Cloud-specific governance (e.g., AWS Control Tower guardrails [2]) constrains portability. We address these gaps with a hybrid evaluation plane coupling rule logic and embedding similarity, backed by an integrity-focused audit pipeline.

II. CONTRIBUTIONS

The work provides: (1) A hybrid evaluation architecture combining deterministic rule filtering and vector similarity; (2) An extensible handler interface enabling semantic and emerging policy types with low change cost; (3) A hash-linked audit log for tamper-evidence without blockchain overhead; (4) A metrics-driven evaluation plan aligning architectural tactics with verifiable quality attributes.

III. SYSTEM ARCHITECTURE

Fig. 1 System Architecture

Core components:

- API Gateway: Authentication, authorization, rate limiting.
- Policy Service: CRUD, versioning, vector indexing (pgvector [4]).
- Evaluation Plane: Cache-first retrieval, rule + semantic evaluation, decision event emission.
- Audit Logger: Consumes decision events; persists hash-chained records.
- Redis: Low-latency hot-set and embedding cache.
- Broker: Decouples evaluation from persistence; supports resilience tests.
- Observability Stack: Metrics (Prometheus [3]), logs, planned traces (OpenTelemetry [5]).




Fig. 1. System Architecture: Hybrid policy evaluation with event-driven decoupling

IV. DESIGN METHODOLOGY

Microservice decomposition isolates policy mutation from evaluation throughput. Event-driven processing reduces synchronous coupling. Vector similarity augments metadata filtering, allowing semantically proximate policy selection. Hash chaining supplies integrity assurance with linear verification.

V. QUALITY ATTRIBUTES

A. Scalability

Tactics: Stateless pods, horizontal autoscaling, Redis caching, asynchronous buffering. Metrics: P95 latency < 150 ms under 2000 requests/sec; cache hit ratio correlated with latency reduction.

B. Extensibility

A PolicyHandler interface:

```
export interface PolicyHandler {
  supports(policyType: string): boolean;
  evaluate(context: any, policy: Policy): Promise<any>;
}
```

New handlers (e.g., advanced semantic constraint) integrate via dependency injection without core modification. Verification: Diff locality; unchanged regression suite; cyclomatic complexity \downarrow 10.

C. Observability & Auditability

Correlation IDs propagate through logs and metrics. Audit integrity uses chained hashes:

```

prev_hash = "GENESIS"
for event in decision_events:
    serialized = canonical_json(event)
    curr_hash = sha256(prev_hash + serialized)
    store({...event, prev_hash, curr_hash})
    prev_hash = curr_hash

```

Periodic verification recomputes sequence and alerts on mismatch; target audit lag \leq 5 s at sustained load.

D. Maintainability

Modular boundaries and standardized evaluation signatures reduce change effort. KPIs: Mean Time To Change for routine policy extension \leq 0.5 developer-day.

VI. IMPLEMENTATION OVERVIEW

Policies stored in PostgreSQL with pgvector embeddings; evaluation pipeline performs: (1) Cache lookup; (2) Vector similarity search (threshold-based); (3) Rule evaluation; (4) Decision event emission; (5) Audit service hash-link persistence.

VII. EVALUATION PLAN

Phases: (1) Functional tests (unit, integration); (2) Load tests (k6 ramp to 2000 req/sec); (3) Resilience (broker partition throttling). Metrics scraped every 5 s; Grafana dashboards visualize throughput, latency, audit lag. Success: No message loss; hash chain validates end-to-end; latency targets met.

VIII. SECURITY CONSIDERATIONS

Current: JWT-based client auth; least-privilege DB roles; rate limiting. Planned: mTLS inter-service; policy signature verification for provenance.

IX. RELATED WORK

OPA [1] offers strong declarative evaluation but lacks semantic retrieval integration and native hash-chained audit logging. AWS Control Tower guardrails [2] provide opinionated cloud governance with limited cross-cloud extensibility. Our approach differentiates through hybrid retrieval, event-driven decoupling, and cryptographic audit integrity.

X. CONCLUSION

The proposed Guardrails-as-a-Service platform aligns architectural tactics with measurable quality attributes. Hybrid policy evaluation and hash-linked auditing together enhance governance fidelity and trust. Future work targets production hardening, expanded semantic handlers, and full distributed tracing.

ACKNOWLEDGMENT

Omitted for anonymity.

REFERENCES

- [1] Open Policy Agent Documentation. [Online]. Available: <https://www.openpolicyagent.org>
- [2] AWS Control Tower Guardrails. [Online]. Available: <https://docs.aws.amazon.com/controlltower>
- [3] Prometheus Documentation. [Online]. Available: <https://prometheus.io>
- [4] pgvector Extension Documentation. [Online]. Available: <https://github.com/pgvector/pgvector>
- [5] OpenTelemetry Specification. [Online]. Available: <https://opentelemetry.io>